Running head: PRAGMATICS OF NEGATION

Negation is only hard to process when it is not relevant or informative

Ann E. Nordmeyer and Michael C. Frank

Department of Psychology, Stanford University

Corresponding author: Ann E. Nordmeyer

Department of Psychology

Stanford University

Building 420 (Jordan Hall)

450 Serra Mall

Stanford, CA 94305

Phone: 650-721-9270

Email: anordmey@stanford.edu

Abstract

Negation is a fundamental element of language and logical systems, but processing negative sentences can be challenging. Early investigations suggested that this difficulty was due to the representational challenge of adding an additional logical element to a proposition, but in more recent work, supportive contexts mitigate the processing costs of negation, suggesting a pragmatic explanation. We make a strong test of this pragmatic hypothesis by directly comparing speakers and listeners. Speakers produce negative sentences more often when they are both relevant and informative. Listeners in turn are fastest to respond to sentences that they expect speakers to produce. Because negative sentences are only difficult in contexts when they are unlikely to be produced, the primary challenges in processing negation are likely pragmatic rather than representational.

Keywords: Language, Psycholinguistics, Language Comprehension, Language Production, Pragmatics

Introduction

Language allows us to describe not only the world as we see it, but also the world as it is not. Nevertheless, for human language users, processing negation is often slow and effortful. Deciding the truth value of a sentence like "star isn't above plus" takes a lot longer than making the same decision about a positive sentence (H. Clark & Chase, 1972; Carpenter & Just, 1975; Just & Carpenter, 1971, 1976). And in language comprehension tasks, participants often show evidence consistent with having processed the positive components of a sentence prior to negating them, suggesting again that negation is challenging (Kaup & Zwaan, 2003; Kaup, Ludtke, & Zwaan, 2006; Hasson & Glucksberg, 2006; Fischler, Bloom, Childers, Roucos, & Perry, 1983; Lüdtke, Friedrich, De Filippis, & Kaup, 2008; Ferguson & Sanford, 2008). Why do adults struggle to process negation despite spontaneously producing negative sentences with ease?

One explanation is that not all negations are equally felicitous. For example, in many contexts it would be somewhat strange to say "my car isn't purple". And it would be even stranger to say "my car doesn't have an ejector seat." On Gricean and neo-Gricean accounts of the pragmatics of language use in context, listeners expect speakers to produce truthful, informative, and relevant utterances (Grice, 1975; Horn, 1984; Levinson, 2000; Sperber & Wilson, 1986). The above utterances are both truthful, but vary in their informativeness and relevance.

In this paper, we define an informative utterance as one that conveys more information about the referent (i.e., makes the referent easier to identify in context; see Frank and Goodman (2012) for a formalization of informativeness, and H. H. Clark (1976) and Moxey (2006) for a discussion of informativeness as it applies to negative sentences). The utterance "my car isn't purple" doesn't convey as much information as, say, "my car

is gray", because "isn't purple" refers to a larger set of cars than "is gray". If we were in a parking lot full of purple cars, however, the statement "my car isn't purple" becomes much more informative, because it helps to uniquely identify the car in question *in that particular context*.

Many formal theories of pragmatics focus specifically on the role of either relevance or informativeness (e.g. Sperber & Wilson, 1986; Frank & Goodman, 2012). In this paper, we attempt to examine the separate effect of these pragmatic factors. We define relevance as a salient or expected feature based on the context. Although under this definition an informative utterance is *always* relevant (because any feature that helps you identify a referent is also going to be something noticeable or salient about that referent in context), a relevant utterance is not necessarily informative. For example, in the above two examples, the sentence "my car isn't purple" is relevant even though it is not usually informative, because color is a salient feature that all cars have (i.e. most people will mention color when asked to describe a particular car, even if the color itself isn't very exciting, e.g., "my car is beige."). The sentence "my car doesn't have an ejector seat", however, is neither relevant nor informative, because in most contexts having an ejector seat is not a feature that we would expect a car to have. If you are James Bond, however, and you live in a universe where ejector seats are sometimes a feature that cars have, then the utterance "my car doesn't have an ejector seat" becomes very relevant (though still not particularly informative, under our definition, because it is still the case that the *majority* of cars don't have ejector seats). Thus, we take relevance and informativeness to be separable pragmatic factors, and both can change depending on the context in which an utterance is produced.

A key element of neo-Gricean theories of communication is that listeners *expect* speakers to produce utterances that are informative and relevant. When an utterance appears to violate these principles, listeners make inferences about the speaker's intended

meaning that go beyond the literal, semantic meaning of the utterance. Neo-Gricean theories have been used to explain the pragmatic inferences that can be drawn from a negative utterance (Moeschler, 1992, see Tian, Breheny, & Ferguson, 2010 for experimental support); for example, if I told you that my car isn't purple, you might *assume* that all of the other cars in the parking lot are purple (because otherwise I would have produced a more informative utterance). When the context rules out this kind of interpretation, however, the utterance is simply infelicitous. Is this kind of pragmatic infelicity generally responsible for the processing cost of negation?

Consistent with this suggestion, presenting negative information in a supportive context can mitigate some of its processing costs (Wason, 1965; Glenberg, Robertson, Jansen, & Johnson-Glenberg, 1999). When a negated feature is explicitly mentioned or inferred in preceding sentences (Lüdtke & Kaup, 2006; Orenes, Beltrán, & Santamaría, 2014), or when negation is presented within a dialogue (Dale & Duran, 2011), negative sentences tend to be processed faster relative to negative sentences presented without context. And in an ERP experiment, contextually-supported negations (e.g., "with proper equipment, scuba-diving isn't very dangerous") elicited smaller N400 responses—a marker of semantic processing costs—than unlicensed negations (e.g., "bulletproof vests aren't very dangerous"; Nieuwland & Kuperberg, 2008).

Although this previous work supports the idea that some kind of contextual expectations are the source of negation's processing cost, they do not specify the precise nature of these expectations. Our current experiment directly tests two hypotheses. First, speakers tend to produce negative sentences only when they are both relevant and informative given the context. Second, expectations about what speakers would likely say—and their match or mismatch with what the speaker in fact *does* say—are responsible for the processing costs of negation. To formalize this second hypothesis, we make use of recent probabilistic models of language comprehension, defining a listener's

pragmatic expectations as the probability that a speaker would utter a statement in order to convey a particular meaning (Frank & Goodman, 2012), and using *surprisal*, an information-theoretic measure of expectation-based processing costs (Levy, 2008), to predict processing times.

The current study

In our experiment, we show participants sets of four characters who varied in terms of appearance (i.e. hair and shirt color) as well as the presence or absence of a target feature (e.g., boys with or without apples). Varying the presence or absence of the target feature created five possible context conditions with a different proportion of characters having the target feature: $\frac{0}{4}$, $\frac{1}{4}$, $\frac{2}{4}$, $\frac{3}{4}$, $\frac{4}{4}$, where the numerator represented the number of characters in the set who e.g., have apples. In the *speaker* condition, we ask participants to produce written descriptions of one of the characters, while in the *listener* condition, we ask participants to evaluate the truth value of negative or positive statements about the same pictures (see Figure 1).

We predict that speakers will be more likely to mention the presence or absence of the feature (e.g., apples) when that feature is relevant (e.g., when at least one other character in the set has apples), and when mentioning that feature is informative (e.g., saying a character "has apples" should be more likely when *few* other characters have apples, and saying that a character "has no apples" should be more likely when *most* other characters have apples). Furthermore, we predict that listeners' processing costs should be predicted by how unlikely it is for a speaker to use the same sentence to describe that picture.

On every trial, a speaker could always produce a relevant and informative utterance by describing the color of the referent's shirt. We were therefore specifically interested in how the context (i.e. the other characters in the set) would influence the probability that

a speaker would mention the presence or absence of the target feature.

Mentioning the target feature is only relevant when at least one other character has the target item; thus, we expected that very few (if any) participants would mention the absence of target items in the $\frac{0}{4}$ context. In the $\frac{1}{4}$ context, at least one character in the set possessed target items, so mentioning the presence or absence of the target item becomes relevant in this context. We therefore predicted a sharp decrease in the surprisal of producing a true negative utterance between the $\frac{0}{4}$ and $\frac{1}{4}$ contexts, because the probability of producing a negative utterances increases from zero (or very close to zero).

Mentioning the absence of target items in the $\frac{1}{4}$ is more informative than mentioning the absence of target items in the $\frac{0}{4}$; however, a negative utterance in this context is still not very informative compared to e.g. describing the color of a character's shirt (i.e., "no apples" describes three possible characters in the $\frac{1}{4}$ condition, whereas "blue shirt" uniquely describes the referent in this context). Therefore, we expected that the probability of producing a negative utterance would increase as the number of context characters with target items increased (i.e., in the $\frac{3}{4}$ context "no apples" is just as informative as "blue shirt", because both utterances uniquely identify the referent). Note that for positive sentences we would predict the opposite effect of context: In the $\frac{4}{4}$ context, saying "has apples" is not very informative, because everyone has apples; in the $\frac{1}{4}$ context, however, "has apples" is very informative because the target character is the only person in the context who has apples.

Method

Participants

We recruited a planned sample of 500 participants to participate in an online experiment through the Amazon's Mechanical Turk (mTurk) website; 5 participants were rejected for indicating that they were under 18 after completing the experiment. Of the

remaining 495 participants, 238 were male and 255 were female, two declined to report gender, and ages ranged from 18 – 65+. We restricted participation to individuals in the US and paid 50 cents for this 10 minute study. Participants were randomly assigned to either the speaker condition (n = 242) or the listener condition (n = 253).

Stimuli

[Figure 1 about here.]

Thirty-two trial items were created in which characters were shown holding either two of the same common, recognizable objects ("target items"; e.g., two apples), or holding nothing. Within each trial, characters' shirt and hair colors also varied, so as to provide other referential possibilities for speakers.

Each participant saw trials in which different proportions of characters were holding target items (context condition). These contexts showed $\frac{0}{4}$, $\frac{1}{4}$, $\frac{2}{4}$, $\frac{3}{4}$, or $\frac{4}{4}$ of the characters holding objects. The order of characters was shuffled on each trial, with the referent of the sentence appearing in a random position.

Participants in the speaker condition saw each image paired with an incomplete sentence (e.g. "[NAME] has ⎯⎯⎯⎯⎯⎯.”). In half of the trials, the highlighted character was holding target items ("item trials"), and in half of the trials, the highlighted character was holding nothing ("nothing trials"). The experiment was fully crossed such that target characters appeared with or without target items (i.e., "referent type") an equal number of times in each context type.

Participants in the listener condition saw the same set of images. On each trial a sentence of the form "[NAME] [has/has no] [TARGET ITEM]" appeared. Half of the sentences were positive and half were negative (sentence type), and they were paired with pictures such that half were true and half were false (truth value), resulting in four possible trial types (true positive, true negative, false positive, and false negative).

Because true positive and false negative sentences cannot occur in a $\frac{0}{4}$ context (i.e. the referent must have the target item in these trials), and true negative and false positive sentences cannot occur in a $\frac{4}{4}$ context, each trial type occurred in four possible contexts. The experiment was fully crossed, with participants receiving eight true positive, eight false positive, eight true negative and eight false negative sentences distributed equally across context types in a randomized order over the course of the study.

Procedure

Participants were first presented with a brief overview screen which explained that they would play a language game. Once participants accepted the task, they were randomly assigned to either the speaker condition or the listener condition, and saw a more detailed instructions screen which explained the task and informed them that they could stop at any time. The speaker and listener conditions of the experiment can be viewed at

`http://anordmey.github.io/negatron/experiments/speakers/negatron.html` and `http://anordmey.github.io/negatron/experiments/listeners/negatron.html`, respectively.

In the speaker condition, participants saw an array of four pictures on each test trial: The target pictures and three context pictures presented in a random horizontal arrangement. Participants were told to look at these pictures for four seconds, at which point a red box appeared around one of the pictures. One second later, an incomplete sentence appeared. Participants were told that the incomplete sentence was about the picture with the box around it, and that their job was to complete the sentence (by typing into a small text box) using only a few words. Participants were explicitly discouraged from referring to the character's position in the array of pictures.

In the listener condition, participants first saw eight positive sentence practice trials

with feedback about incorrect responses before beginning the test trials. In each test trial, participants saw an array of four pictures presented in a random horizontal arrangement. Participants were told to look at these pictures for four seconds, at which point a red box appeared around one of the pictures. One second later, a sentence about that picture appeared. Participants were told to read the sentence and respond as quickly and accurately as possible with a judgment of whether it was true or false when applied to the highlighted picture. We recorded reaction times for each trial, measured as the time from when the sentence was presented to the moment when the response was made.

Data Processing

We excluded 16 participants who did not list English as their native language and two participants from the listener condition for having an overall accuracy below 80% (overall accuracy on this task was very high, with participants responding correctly on 98% of true positive trials, 93% of true negative trials, 96% of false positive trials, and 97% of false negative trials). This left a total of 477 participants for analysis (244 in listener condition, 233 in speaker condition).

In the speaker condition, we coded participant's productions. Affirmative responses labeling the target feature were coded as "positive" (e.g., "apples," "two apples," "red apples," etc.). Responses negating the target feature (e.g., "no apples") were coded as "negative." All other responses (e.g. descriptions of the characters' clothing or hair color, as well as other types of positive or negative utterances) were coded as "other.". Codes were hand-checked to ensure that label synonyms or spelling errors were coded correctly.

In the listener condition, we excluded trials with RTs greater than 3 standard deviations from the log-transformed mean, a criterion established in our previous experiments (Nordmeyer & Frank, 2014).

To predict responding in the listener condition, we used productions from the

speaker condition. We calculated the proportion of positive sentences describing characters who possessed target items, and the proportion of negative sentences describing characters with nothing, creating probability distributions for true positive and true negative utterances in each context. We then used this distribution to calculate the surprisal of hearing a true positive or true negative sentence for each context. Surprisal is an information-theoretic measure of the amount of information carried by an event (in this case, the amount of information conveyed by a sentence); in prior work on sentence comprehension it has been used successfully as a linking hypothesis between production probabilities and reaction times (Levy, 2008). Surprisal (or "self-information" I) for a sentence s is defined as

$$I(s) = -\log(P(s)). \tag{1}$$

Results

Speaker Condition

[Figure 2 about here.]

[Figure 3 about here.]

Participants were more likely to produce true positive utterances about target items than they were to produce true negative utterances about the *absence* of target items. As the number of characters in the context with target items increased, however, participants became less likely to produce true positive utterances about referents with target items, and less likely to produce true negative utterances about characters without target items (see Figure 2). These findings support our hypothesis that speakers will produce informative utterances (i.e. produce sentences that are maximally effective at identifying the referent), because true positive utterances about the presence of target items are more

informative when none of the other characters have target items (e.g. the $\frac{1}{4}$ context), whereas true negative utterances about the absence of target items are more informative when the other characters *do* have target items (e.g. the $\frac{3}{4}$ context)

Consistent with our hypothesis that speakers will produce negative sentences that are relevant (i.e. negating a salient feature of the context), negative sentences were almost never produced on nothing trials in the $\frac{0}{4}$ context. That is, "no apples" was a very rare production on trials where there weren't any apples in the context; participants in this condition tended to produce affirmative sentences describing other features of the referent, such as their clothing or hair color. Participants only consistently start producing negative utterances in the $\frac{1}{4}$ context, where at least one other character has target items, thus making the target items a relevant feature to discuss.

To evaluate the reliability of these patterns, we fit two separate binomial mixed-effects models to test the probability of producing positive utterances about target items on item trials, and the probability of producing negative utterances about target items on nothing trials. To test the quantitative effects of context, we coded this variable as numeric, e.g. the proportion of characters in the context with target items. In addition, to test the effect of relevance on producing true negative utterances, we created a dummy code to separately test for the effects of the $\frac{0}{4}$ context compared to all of the other contexts. All mixed-effects models used the maximal convergent random effects structure and were fit using the lme4 package version 1.1-7 in R version 3.1.2. Raw data and analysis code can be found at `https://github.com/anordmey/negatron`.

To test the effect of context on the probability of producing a true negative utterance (i.e., negative utterances describing the absence of target items on nothing trials), we fit a model with the following specification: `negation ~ dummy context + numeric context + (1 | subject) + (1 | item)`. We found a significant effect of the dummy coded context, indicating that participants were less likely to produce negative

sentences in the $\frac{0}{4}$ context compared to the other contexts ($\beta = -1.19$, $p = .015$), suggesting a significant effect of relevance on the production of true negative sentences. We also found a significant positive linear effect of context, with the probability of producing a negative sentence increasing as the proportion of characters with target items increases ($\beta = 7.30$, $p < .001$). A model predicting the effect of context on the probability of producing a true positive utterance (i.e., a sentence describing the presence of target items on item trials; model specification `positive` $\sim$ `numeric context +` `(1 | subject) + (1 | item)`) found a significant *negative* linear effect of context ($\beta = -5.5$, $p < .001$), indicating that the probability of producing a positive utterance decreases as the proportion of characters with target items increases. These findings indicate that the probability of producing a true positive or a true negative utterance is influenced by the informativeness of that utterance in context.

Figure 3 shows the surprisal of true positive and true negative sentences. Consistent with our predictions, surprisal is highest for true negative sentences in the $\frac{0}{4}$ context, indicating an effect of relevance. Surprisal for decreases for true negative sentences and increases for true positive utterances as the number of context characters with target items increases, indicating a significant effect of informativeness on the production of both positive and negative sentences. The steep gap in surprisal between the $\frac{0}{4}$ and the $\frac{1}{4}$ contexts is driven by the fact that in the $\frac{0}{4}$ context, the probability of producing a negative utterance is close to zero; therefore, the surprisal for these utterances is much higher than the surprisal of an utterance that is even slightly more probable.

In the above analysis (and in the "Predicting Listeners with Speakers" section, below), we coded utterances such as "Bob has nothing" as "other" rather than "negative." We made this decision because we wanted to examine speaker utterances that were as close as possible to the kinds of utterances that we presented to participants in the listener condition. Furthermore, our definition of relevance is specific to the negation of the target

item (e.g., negating "apples", not just producing any negation). Consistent with this prediction, if we include other types of negation in our analysis (e.g., any utterance with "not", "without", "zero", etc.), the results of the speaker condition are the same as those reported above (i.e., significant effects of relevance and informativeness). If we include instances of "nothing", however, the effect of informativeness remains highly significant ($\beta = 6.55$, $p < .001$), but the effect of relevance is no longer significant ($\beta = -0.21$, $p = 0.50$). That is, "has no apples" is not a very relevant utterance in a context where no one else has apples, but "has nothing" is, perhaps, a relevant thing to say in the context of an experiment where people sometimes have objects and sometimes do not.

Listener Condition

[Figure 4 about here.]

Participants were fastest to respond to true positive sentences, and slowest to respond to true negative sentences, replicating previous findings (H. Clark & Chase, 1972). Listeners' responses to true negative sentences mirrored the surprisal of true negative sentences in the speaker condition, with participants responding slowest to true negatives in the $\frac{0}{4}$ context. The same pattern was seen in response to false positive sentences, suggesting that listeners expect speakers to describe relevant features of the context even when the sentence is false (see Figure 4).

We fit a linear mixed-effects model to examine the interaction between sentence type (positive or negative), truth value (true or false), and context as predictors of reaction time. The model specification was as follows: `RT ~ sentence × truth × context + (sentence | subject) + (sentence | item)`. All model coefficients are shown in Table 1. In addition to main effects of sentence type ($\beta = -205$, $p < .001$) and truth value ($\beta = -384$, $p < .001$), the model showed an interaction such that true positive sentences elicited the fastest responses and true negative sentences elicited the slowest responses

($\beta = -663$, $p < .001$). The model showed a significant negative linear effect of context, with reaction times decreasing as the proportion of characters with target items increased ($\beta = -341$, $p < .001$). A significant three-way interaction between sentence type, truth value, and context indicates that this pattern was driven primarily by responses to true negative sentences, however, with context having the most pronounced effect on true negative utterances ($\beta = -901$, $p < .001$).

To explore the separate effects of relevance (i.e., the effect of the $\frac{0}{4}$ context compared to the others) and informativeness (i.e. the linear effect of context) on responses to true utterances, we fit two separate models to reaction times in response to true positive and true negative utterances (model specification for true negative model: `rt ~ dummy context + numeric context + (dummy context + numeric context | subject) + (dummy context + numeric context | item)`; model specification for true positive model: `rt ~ numeric context + (numeric context | subject) + (numeric context | item)`. We found a significant effect of the dummy coded context variable on reaction times for negative utterances, with the $\frac{0}{4}$ context producing significantly slower reaction times compared to the other contexts ($\beta = 307$, $p < .001$). We did not find a significant linear effect of context above and beyond the effect of the $\frac{0}{4}$ context ($\beta = -41$, $p = .58$). We did, however, find a significant positive linear effect of context on the the reaction time to respond to positive sentences ($\beta = 111$, $p = .02$), indicating a significant effect of informativeness on RTs to positive sentences, but not negative sentences.

[Table 1 about here.]

Predicting Listeners with Speakers

[Figure 5 about here.]

To test the hypothesis that processing times are a function of listeners' expectations about what a speaker will say, we regressed the mean reaction time in response to true

positive and negative utterances in each condition against the surprisal for the same utterances (Figure 5). There was a significant positive relationship between surprisal and reaction time for true negative sentences, $r^2 = .93$, $p < .001$, supporting our prediction that the effects of context on reaction time reflect differences in how speakers would describe the same stimuli. This relationship between surprisal and reaction time for true negative sentences holds even with the outlying $\frac{0}{4}$ context removed from analysis ($r^2 = .71$, $p = .017$).

In real world contexts, there are many possible utterances that a speaker might produce in a given moment. How is this relationship between speaker surprisal and listener reaction time influenced by the availability of other possible utterances? In a previous, unpublished version of this study, participants engaged in a task that was identical except that all four characters were identical except for the presence or absence of target items. Thus, in the previous version, true utterances such as "has a blue shirt" were always relatively uninformative (because they were true of all characters in a given trial). A negative utterance on nothing trials was therefore equally or more informative compared to describing shirt color in the $\frac{1}{4}$, $\frac{2}{4}$, and $\frac{3}{4}$ contexts. Contrast this with the current study, where describing the color of a characters shirt is *always* an informative utterance; the informativeness of a true negative utterance is only comparable to the these competing utterances in the $\frac{3}{4}$ context. The effect of the informativeness of these competing utterances is reflected in the raw probabilities for the speaker data (Figure 2), where there is a jump in the probability of producing a true negative utterance between the $\frac{2}{4}$ and the $\frac{3}{4}$ context in this version of the experiment. In the previous version, this sharp increase occurred between the $\frac{0}{4}$ and $\frac{1}{4}$ contexts – again, the point where producing a true positive utterance becomes as informative as producing an utterance about e.g., shirt color. In both studies, however, the *surprisal* of producing a negative utterance is significantly higher in the $\frac{0}{4}$ than in any other context, because the probability of

producing a true negative utterance in this context is close to zero. In both the current version and the previous version of these experiments, in the listener condition the $\frac{0}{4}$ context yields much slower reaction times than any other context. Critically, replicating both the direction and magnitude of the findings reported here, in the previous study, speaker surprisal and lister reaction time were strongly related, $r^2 = .89$, $p < .001$. Thus, the relationship between speaker surprisal and listener reaction time appears to be robust to changes in the relative informativeness of competing utterances.

General Discussion

What makes negation hard to process? While previous work has proposed that processing negative elements is especially difficult because of features intrinsic to negation, our work here suggests that general pragmatic mechanisms are responsible for much of the difficulty associated with negation. Negative sentences presented without context are uninformative and irrelevant; thus, they are unlikely to be produced by speakers. In turn, listeners respond to these unlikely utterances with increased processing times. In contexts where negation is more informative and relevant, processing costs are lower. Overall, this evidence supports a Gricean interpretation of the processing of negation.

While previous work has shown that contextual factors facilitate the processing of negation (Wason, 1965; Nieuwland & Kuperberg, 2008; Dale & Duran, 2011; Orenes et al., 2014), our findings here go further. First, by using actual language productions as the predictor of processing difficulty, our work strongly implicates specifically pragmatic factors. Because the field of pragmatics is concerned with language *use*, demonstrating a relationship between actual speaker productions and listener processing time is critical to the argument that pragmatic factors are responsible for the processing cost of negation, and to our knowledge none of the past work on the effect of context on negation has demonstrated this relationship. Second, rather than treating pragmatics as a black box,

we show that two different components—informativeness and relevance—each contribute to the relative (un-)likelihood of hearing a negation. Speakers were unlikely to mention the absence of a feature unless it was a relevant feature in the context, and were more likely to mention the presence or absence of features when mentioning those features would help uniquely identify the referent (i.e., when their utterances were informative). The fact that informativeness played a role in the production of both positive and negative utterances supports our argument that these pragmatic pressures are general, rather than specific to negation.

Listeners' reaction times were highly correlated with the surprisal of a speaker producing the same utterances in context. For true negative sentences, however, this effect was driven primarily by the effect of relevance (i.e., the difference between the $\frac{0}{4}$ and all other contexts). Why didn't informativeness play a role in listeners' response times? One possibility is that only relevance plays a role in forming listeners' expectations about what a speaker will say. Another possibility is that the effect of informativeness (for both listeners and speakers) is small compared to the effect of relevance, and that the reaction time measure is too noisy to identify this effect. Although our results cannot disentangle these two possibilities, the fact that we did find a significant effect of informativeness on listeners' responses to true positive sentences suggests that informativeness does play some role in listeners' expectations.

Our work here also uses surprisal, an information-theoretic measure of processing difficulty, as the linking hypothesis between speaker probabilities and reaction times (Levy, 2008). Although this link has substantial support in the realm of syntactic processing (Demberg & Keller, 2008; Boston, Hale, Kliegl, Patil, & Vasishth, 2008), to our knowledge, our findings are the first example of using surprisal over sentence-level pragmatic expectations, rather than word-level syntactic expectations. This success suggests that, in concert with the appropriate predictive models, surprisal theory could be

productively applied to the prediction of processing difficulty beyond the level of syntax.

Although our focus here was on negation, our findings have implications for sentence processing more generally. Debates about the effects of pragmatics on linguistic processing exist in other domains, such as the processing of scalar implicatures (the pragmatic inference that e.g., "some" implicates "some but not all"; Huang & Snedeker, 2009, 2011; Grodner, Klein, Carbary, & Tanenhaus, 2010). Tomlinson, Bailey, and Bott (2013) provide an informative comparison between scalar implicature and negation, presenting mouse-tracking trajectories for each. Their negation data show the same pattern of processing difficulties we observe, and critically, their data on the processing of underinformative "some" utterances look almost identical. We hypothesize that, in both cases, participants' processing difficulty is a function of the violation of their pragmatic expectations.

In sum, our findings here suggest that a large part of the processing difficulties of negative sentences arise from the relative pragmatic felicity of negation in context. They do not rule out the possibility that there is some cost to processing an additional logical element, but this processing cost would have to be quite small with respect to the magnitude of the effects we observed (and previous measurements). This finding leads us to the following conclusion: When logical words are used in a communicative context, we have no difficulty understanding them.

Author Contributions

Both authors developed the study concept and contributed to the study design. Data collection was conducted by A. E. Nordmeyer. A. E. Nordmeyer performed the data analysis and interpretation under the supervision of M.C. Frank. Both authors contributed to the development of the manuscript and approved the final version of the manuscript for submission.

References

Boston, M., Hale, J., Kliegl, R., Patil, U., & Vasishth, S. (2008). Parsing costs as predictors of reading difficulty: An evaluation using the potsdam sentence corpus. *The Mind Research Repository (beta)*.

Carpenter, P., & Just, M. (1975). Sentence comprehension: A psycholinguistic processing model of verification. *Psychological Review*, *82*, 45–73.

Clark, H., & Chase, W. (1972). On the process of comparing sentences against pictures. *Cognitive Psychology*, *3*, 472–517.

Clark, H. H. (1976). Semantics and comprehension.

Dale, R., & Duran, N. (2011). The cognitive dynamics of negated sentence verification. *Cognitive Science*, *35*, 983–996.

Demberg, V., & Keller, F. (2008). Data from eye-tracking corpora as evidence for theories of syntactic processing complexity. *Cognition*, *109*, 193–210.

Ferguson, H., & Sanford, A. (2008). Anomalies in real and counterfactual worlds: An eye-movement investigation. *Journal of Memory and Language*, *58*, 609-626.

Fischler, I., Bloom, P., Childers, D., Roucos, S., & Perry, N. (1983). Brain potentials related to stages of sentence verification. *Psychophysiology*, *20*, 400–409.

Frank, M., & Goodman, N. (2012). Predicting pragmatic reasoning in language games. *Science*, *336*, 998.

Glenberg, A., Robertson, D., Jansen, J., & Johnson-Glenberg, M. (1999). Not propositions. *Journal of Cognitive Systems Research*, *1*, 19–33.

Grice, H. (1975). Logic and conversation. *1975*, 41–58.

Grodner, D. J., Klein, N. M., Carbary, K. M., & Tanenhaus, M. K. (2010). Some, and

possibly all, scalar inferences are not delayed: Evidence for immediate pragmatic enrichment. *Cognition*, *116*, 42–55.

Hasson, U., & Glucksberg, S. (2006). Does understanding negation entail affirmation? An examination of negated metaphors. *Journal of Pragmatics*, *38*, 1015–1032.

Horn, L. R. (1984). Toward a new taxonomy for pragmatic inference: Q-based and R-based implicature. In D. Schiffrin (Ed.), *Meaning, form, and use in context: Linguistic applications* (pp. 11–42). Washington, D.C.: Georgetown University Press.

Huang, Y. T., & Snedeker, J. (2009). Online interpretation of scalar quantifiers: Insight into the semantics–pragmatics interface. *Cognitive psychology*, *58*, 376–415.

Huang, Y. T., & Snedeker, J. (2011). Logic and conversation revisited: Evidence for a division between semantic and pragmatic content in real-time language comprehension. *Language and Cognitive Processes*, *26*, 1161–1172.

Just, M., & Carpenter, P. (1971). Comprehension of negation with quantification. *Journal of Verbal Learning and Verbal Behavior*, *10*, 244–253.

Just, M., & Carpenter, P. (1976). Eye fixations and cognitive processes. *Cognitive Psychology*, *8*, 441–480.

Kaup, B., Ludtke, J., & Zwaan, R. (2006). Processing negated sentences with contradictory predicates: Is a door that is not open mentally closed? *Journal of Pragmatics*, *38*, 1033–1050.

Kaup, B., & Zwaan, R. (2003). Effects of negation and situational presence on the accessibility of text information. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, *29*, 439–446.

Levinson, S. C. (2000). *Presumptive meanings: The theory of generalized conversational implicature.* Cambridge, MA: MIT Press.

Levy, R. (2008). Expectation-based syntactic comprehension. *Cognition*, *106*, 1126–1177.

Lüdtke, J., Friedrich, C., De Filippis, M., & Kaup, B. (2008). Event-related potential

correlates of negation in a sentence-picture verification paradigm. *The Journal of Cognitive Neuroscience*, *20*, 1355–1370.

Lüdtke, J., & Kaup, B. (2006). Context effects when reading negative and affirmative sentences. In *Proceedings of the 28th Annual Conference of the Cognitive Science Society* (pp. 1735–1740).

Moeschler, J. (1992). The pragmatic aspects of linguistic negation: Speech act, argumentation and pragmatic inference. *Argumentation*, *6*, 51–76.

Moxey, L. M. (2006). Effects of what is expected on the focussing properties of quantifiers: A test of the presupposition-denial account. *Journal of Memory and Language*, *55*, 422–439.

Nieuwland, M., & Kuperberg, G. (2008). When the truth is not too hard to handle. *Psychological Science*, *19*, 1213.

Nordmeyer, A. E., & Frank, M. C. (2014). A pragmatic account of the processing of negative sentences. In *Proceedings of the 36th Annual Meeting of the Cognitive Science Society*.

Orenes, I., Beltrán, D., & Santamaría, C. (2014). How negation is understood: Evidence from the visual world paradigm. *Journal of Memory and Language*, *74*, 36–45.

Sperber, D., & Wilson, D. (1986). *Relevance: Communication and cognition*. Oxford, UK: Blackwell Publishing.

Tian, Y., Breheny, R., & Ferguson, H. J. (2010). Why we simulate negated information: A dynamic pragmatic account. *The Quarterly Journal of Experimental Psychology*, *63*, 2305–2312.

Tomlinson, J. M., Bailey, T. M., & Bott, L. (2013). Possibly all of that and then some: Scalar implicatures are understood in two steps. *Journal of Memory and Language*, *69*, 18–35.

Wason, P. (1965). The contexts of plausible denial. *Journal of Verbal Learning and Verbal*

*Behavior*, *4*, 7–11.

List of Figures

Figure 1. An example of a true negative trial with a 0/4 context (left) and a 3/4 context (right). The sentence "Bob has no apples" in the 0/4 context is both uninformative (because the sentence is true of all of the characters) and irrelevant (because apples are not present in the context), whereas the same sentence in the 3/4 context is both informative and relevant.

Figure 2. Probability of producing negative sentences on nothing trials (i.e., true negatives) and positive sentences on item trials (i.e., true positives) across different contexts. Negative sentences are shown in grey, and positive sentences in black. The context is notated by a fraction representing the number of characters in the context who held target items. Error bars show 95% confidence intervals computed by non-parametric bootstrapping.

Figure 3. Surprisal for true positive and true negative sentences across different contexts. Negative sentences are shown in grey, and positive sentences in black. The context is notated by a fraction representing the number of characters in the context who held target items. Error bars show 95% confidence intervals computed by non-parametric bootstrapping.

Figure 4.   Reaction times for each trial type across different conditions. Responses to true sentences are shown on the left, and false sentences are shown on the right. Negative sentences are shown in grey, and positive sentences in black. The context is notated by a fraction representing the number of characters in the context who held target items. Error bars show 95% confidence intervals computed by non-parametric bootstrap.

Figure 5.   Reaction times in the listener condition plotted by surprisal in the speaker condition; each point represents a measurement for sentence type and context. Negative sentences are shown in grey, and positive sentences in black. The light gray band shows the linear regression and 95% confidence region for all conditions; the dark gray band shows the linear regression and 95% confidence region excluding the outlying 0/4 condition. Error bars on the horizontal and vertical axes represent 95% confidence intervals on their respective measures.

List of Tables

Table 1

Coefficient estimates from a mixed-effects model predicting listeners' reaction times in response to sentences in different contexts.

	Coefficient	Std. err.	t
Intercept	1598	41	39.16
Sentence (Negative)	-205	38	-5.36
Truth (True)	-384	36	-10.57
Context	-341	42	-8.14
Sentence $\times$ Truth	663	54	12.16
Sentence $\times$ Context	377	59	6.40
Truth $\times$ Context	454	59	7.71
Sentence $\times$ Truth $\times$ Context	-901	83	-10.79